

How does an operating system boot?

- firstly we need to understand the minimum details as to how the IBM-PC boots
- secondly we need to understand the desired final outcome at the end of the boot phase
- thirdly we can examine a specific example to better understand the steps taken to achieve the outcome

How does an IBM-PC boot?

- the bios settings dictates the boot device order
- the bios attempts to load in the first 512 bytes (boot sector) from the various devices in order
- not all devices may be present: usb memory stick, floppy disk
- the bios loads in the 512 bytes from the first found device at 0x7c00
 - it sets register: `cl` to the device number
 - the bios then jumps to location 0x7c00

Boot sector code characteristics and functionality

- it remembers the bios boot device (usb/floppy/harddisk) in a processor register `dl`
- it copies itself into a sensible location (typically out of the way in high memory)
- it reassigns the stack to a consistent location
- it loads in the secondary boot stage of the operating system (a sector at a time)

Boot sector code characteristics and functionality

- it may perform very limited checking as each subsequent sector/track is read from the device into memory
- finally it jumps to the start of the secondary code.

The language used to implement the bootsector (`first`)

- the boot sector (`first`) is not normally written in a high level language as it needs the ability to:
 - copy its code segment, reassign the stack (change the `SP` and stack segment registers)
 - the ability to jump to a physical location and it must fit in 512 bytes

Final desired outcome after all the boot phases are complete



Example: LuK booting

- LuK consists of a collection of modules
- the microkernel only links the modules actually required at runtime
- the mixture of the modules required for different targets and applications may be different

Linker

- uses the file `init` to generate a list of modules and generates an ELF 32 bit x86 executable which contains data, code and symbol information
 - for example we will name this, *application.third*

Boot phases

- in the build directory you would see
 - first, second

first

- first is a tiny model 8086 executable, written in assembly language
- see `luk-1.0.3/boot/BAS/boot.S`
 - watch out as the assembler uses: `mov dest, src`
- its total size (data + code) must not exceed 512 bytes
- its duty is threefold
 - pretend to be a fat12 file system!
 - move itself to a sane location
 - load in `second`

second

- is written in Modula-2, which is compiled and linked into a tiny model 8086 executable
- tiny model
 - sets all segment registers to the same value
 - total size of data + code + stack must not exceed 64K
- in fact due to legacy booting via the floppy disk it cannot be more than 7K

second

- its duty is to load in the *application.third*
- set up protected mode and move from tiny model into 32 bits
- pass various system parameters into *application.third*
 - such as memory size, video memory start
- finally jump to the start of *application.third*

Goal of the overall boot procedure



Goal of the overall boot procedure

- notice that no tiny model code will exist in the end
- all code is 32 bit and belongs to the core microkernel
- `first` and `second` will be overwritten

Overview of the boot stages

- three boot phases
 - *first* boot stage (boot sector, 1 sector, assembly language)
 - *second* boot stage (up to 14 sectors 8088 small mode Modula-2)
 - *LuK* (up to 512K of 32 bit code, Modula-2 and C)



LuK boot first

- (program first)
- 512 bytes boot sector is small! Just enough space to place an assembly language program which loads in a larger program
 - loads in *secondary* boot stage at 0x90200
 - jumps to 0x90200
- *secondary* boot stage (program second)
 - consists of limited amounts of assembly language
 - most of the code is written in Modula-2 but compiled to small mode 8088
 - the secondary stage may be up to 14 sectors in size (14 * 512 bytes)

Secondary boot stage

- purpose of *secondary* boot stage is to load in your *application.third* code as quickly as possible
 - it uses whole track reads whenever possible (fast)
 - the *primary* boot stage only used single sector loads (slow)
 - it loads in the LuK 32 bit executable (*application.third*) into location 0x10000
 - collects vital statistics about the PC (how much memory the PC contains and where video memory starts)
 - saves this information
 - turns the floppy disk motor off

- finally *second* puts the microprocessor into 32 bit mode and calls *application.third*

Boot phase in more detail

- *how* do you put LuK in the right place?
 - tip, think backwards

- start with the final position you desire
 - and consider how you can achieve it
 - draw memory maps of the different LuK bootstage intermediate positions

Final memory map for LuK



Second memory map for LuK



Boot memory map for LuK



Conclusion

- this technique works
- it is not the most efficient, it might be possible to make first perform the actions of second
- however the approach presented here allows us to:
 - execute high level language code sooner
- some of the older limits should be removed now that booting floppy disks is no longer needed
- maybe it would be sensible to move LuK to start at 1MB upwards
 - would allow LuK to expand