

Taxonomy of client/server architectures

- so far we have looked at a simple TCP server/client and a simple UDP server/client
- this week we will further classify these servers
- first we will examine the pros and cons of the TCP and UDP server/clients

The pros/cons for a TCP client server

- pro connection is reliable
- pro reasonably efficient for sending medium/large amounts of data
- con requires packets to be sent (overhead) to setup the connection and close the connection
- con inefficient to send tiny amounts of data

The pros/cons for a UDP client server

- pro simpler than the TCP counterpart
- pro very efficient for sending tiny amounts of data
- pro no connection is created by UDP, hence less overhead
- con it uses the UDP transport thus data might be scrambled or lost in transit
 - connectionless transport characteristics
- con you have to manage the unreliable nature of the connection yourself
 - examples NFS, VoIP

Returning to the basic server algorithm for TCP or UDP

- conceptually each server follows a simple algorithm, expressed in pseudo code:

- ```
it creates a socket
binds the socket to a well known port
loop
 accept the next client
 request from this port
 serve this request
 formulate a reply
 send the reply to client
end
```

## Problems with the simple server?

- unfortunately this is only good enough for simple applications
- consider a service requiring considerable time to handle each request
  - example suppose a file transfer client server were implemented like this!
  - one user requests a huge file
  - moments later another user might wish to transfer a small file

## Problems with the simple server?

- the second user has to wait a considerable time just to transfer a small file
- the second user is **blocked** until the first user has finished with the server
- thus servers are seldom built like this

## Taxonomy of client/server architecture

- first on the list in our taxonomy of client servers is
- **iterative server** (as we have just seen)
  - used to describe a server implementation that processes one request at a time

## Taxonomy of client/server architecture

- second on the list in our taxonomy of client servers is a
  - **concurrent server**
    - used to describe a server that handles multiple requests at a time
- best viewed from the client perspective
  - the server appears to communicate with multiple clients concurrently.
- *the term concurrent server refers to whether the server handles multiple requests concurrently, not to whether the underlying implementation uses multiple concurrent processes*



## Concurrent server pro/cons

- concurrent servers are more difficult to design and build
  - the resulting code is more complex
  - difficult to modify
- most programmers choose concurrent server implementations

## Iterative server pro/cons

- cause unnecessary delays in distributed applications
- may be a performance bottleneck that effects many client applications
- *iterative server implementations, which are easier to build and understand, may result in poor performance because they make clients wait for service. Whereas in contrast, concurrent server implementations, which are more difficult to build, yield better performance.*

## Iterative server pro/cons

- we can view these two categories across the TCP/UDP division below:

box with .sw at (0.787,7.941) width 2.165 height 1.378 box with .sw at (0.787,6.563) width 2.165 height 1.378 box with .sw at (2.953,6.563) width 2.165 height 1.378 box with .sw at (2.953,7.941) width 2.165 height 1.378 "concurrent" at 1.181,7.581 ljust "concurrent" at 3.543,7.581 ljust "connectionless" at 1.181,7.187 ljust "connection oriented" at 3.543,7.187 ljust "(UDP)" at 2.362,7.187 ljust "(TCP)" at 3.543,6.793 ljust "(TCP)" at 3.543,8.171 ljust "iterative" at 1.181,8.959 ljust "iterative" at 3.543,8.959 ljust "connection oriented" at 3.543,8.565 ljust "connectionless" at 1.181,8.565 ljust "(UDP)" at 2.362,8.565 ljust

## Pseudo code for the iterative connectionless server



```
create a socket and bind
 to a well known address
 for which a service is
 being offered

loop
 read next request from client
 process the request
 send reply back to client
end
```


## Pseudo code for the concurrent connectionless server

```
create a socket and bind
 to the well known address
 for the service being offered

leave the socket unconnected

loop
 call recvfrom to obtain the
 next client request
 if (fork() == 0) {
 /* child process. */
 process the request
 form a reply and send
 it to client
 (use sendto)
 exit (0)
 }
 /* only the parent gets here. */
end
```

## Pseudo code for a concurrent connection oriented server



```
create a socket and bind
 it to the well known address
 for the service being offered

place socket into passive mode
 making it ready for use by
 the server
```

# Pseudo code for a concurrent connection oriented server

```
loop
 call accept to receive the
 next request from a client
 if (fork() == 0) {
 /* must be the child */
 repeat
 read request from client
 process the request
 form a reply and send
 it to client
 until client wishes to quit
 close connection
 exit (0)
 }
 /* only the parent gets here. */
end
```

## When to use each server type

- iterative vs concurrent
  - iterative server is easier to design, implement and maintain
  - concurrent server can provide a quicker response to requests
  
- use iterative implementation if
  - the time to **process the request** is small



## When to use each server type

- connection oriented vs connectionless
  - connection oriented access means using TCP
    - implies reliable delivery
  - because connectionless transport means using UDP
    - it implies unreliable delivery

## Conclusion

- only use connectionless transport if the application protocol handles reliability
  - or the local area network exhibits:
    - low packet loss
    - no packet reordering (very few do)
- use connection oriented transport whenever
  - a wide area network separates client and server
- never move a connectionless client and server to a wide area network
  - without checking to see if the application protocol handles the reliability problems