Taxonomy of client/server architectures

- so far we have looked at a simple TCP server/client and a simple UDP server/client
- this week we will further classify these servers
- first we will examine the pros and cons of the TCP and UDP server/clients

The pros/cons for a TCP client server

- pro connection is reliable
- pro reasonably efficient for sending medium/large amounts of data
- con requires packets to be sent (overhead) to setup the connection and close the connection
- con inefficient to send tiny amounts of data

The pros/cons for a UDP client server

- pro simpler than the TCP counterpart
- pro very efficient for sending tiny amounts of data
- pro no connection is created by UDP, hence less overhead
- con it uses the UDP transport thus data might be scrambled or lost in transit
  - connectionless transport characteristics
- con you have to manage the unreliable nature of the connection yourself
  - examples NFS, VoIP

Returning to the basic server algorithm for TCP or UDP

- conceptually each server follows a simple algorithm, expressed in pseudo code:

```
it creates a socket
binds the socket to a well known port
loop
   accept the next client
       request from this port
   serve this request
   formulate a reply
   send the reply to client
end
```

Problems with the simple server?

- unfortunately this is only good enough for simple applications

- consider a service requiring considerable time to handle each request
  - example suppose a file transfer client server were implemented like this!
  - one user requests a huge file
  - moments later another user might wish to transfer a small file

Problems with the simple server?

- the second user has to wait a considerable time just to transfer a small file
- the second user is **blocked** until the first user has finished with the server
- thus servers are seldom built like this

Taxonomy of client/server architecture

- first on the list in our taxonomy of client servers is

- **iterative server** (as we have just seen)
  - used to describe a server implementation that processes one request at a time

Taxonomy of client/server architecture

- second on the list in our taxonomy of client servers is a

- **concurrent server**
  - used to describe a server that handles multiple requests at a time

- best viewed from the client perspective
  - the server appears to communicate with multiple clients concurrently.

- *the term concurrent server refers to whether the server handles multiple requests concurrently, not to whether the underlying implementation uses multiple concurrent processes*

Concurrent server pro/cons

- concurrent servers are more difficult to design and build
  - the resulting code is more complex
  - difficult to modify

- most programmers choose concurrent server implementations

Iterative server pro/cons

- cause unnecessary delays in distributed applications
- may be a performance bottleneck that effects many client applications
- *iterative server implementations, which are easier to build and understand, may result in poor performance because they make clients wait for service. Whereas in contrast, concurrent server implementations, which are more difficult to build, yield better performance.*

Iterative server pro/cons

- we can view these two categories across the TCP/UDP division below:

iterative connectionless (UDP)	iterative connection oriented (TCP)
concurrent connectionless (UDP)	concurrent connection oriented (TCP)

Pseudo code for the iterative connectionless server

```
create a socket and bind
   to a well known address
   for which a service is
   being offered

loop
   read next request from client
   process the request
   send reply back to client
end
```

Pseudo code for the concurrent connectionless server

```
create a socket and bind
    to the well known address
    for the service being offered

leave the socket unconnected

loop
   call recvfrom to obtain the
        next client request
   if (fork() == 0) {
       /* child process.  */
       process the request
       form a reply and send
           it to client
       (use sendto)
       exit (0)
   }
   /* only the parent gets here.  */
end
```

Pseudo code for a concurrent connection oriented server

```
create a socket and bind
   it to the well known address
   for the service being offered

place socket into passive mode
   making it ready for use by
   the server
```

Pseudo code for a concurrent connection oriented server

```
loop
   call accept to receive the
   next request from a client
   if (fork() == 0) {
       /* must be the child */
       repeat
          read request from client
          process the request
          form a reply and send
             it to client
       until client wishes to quit
       close connection
       exit (0)
   }
   /* only the parent gets here.  */
end
```

When to use each server type

- iterative vs concurrent
  - iterative server is easier to design, implement and maintain
  - concurrent server can provide a quicker response to requests

- use iterative implementation if
  - the time to process the request is small

When to use each server type

- connection oriented vs connectionless
  - connection oriented access means using TCP
    - implies reliable delivery
  - because connectionless transport means using UDP
    - it implies unreliable delivery

Conclusion

- only use connectionless transport if the application protocol handles reliability
  - or the local area network exhibits:
    - low packet loss
    - no packet reordering (very few do)

- use connection oriented transport whenever
  - a wide area network separates client and server

- never move a connectionless client and server to a wide area network
  - without checking to see if the application protocol handles the reliability problems